

MapReduce

**A programming model
for processing large datasets**

Prof. Tahar Kechadi

School of Computer Science & Informatics

Outline

- **MapReduce, the model**
- **Hadoop, an open-source implementation**



- **Pig, a high-level abstraction layer**



What is MapReduce?

- A data-parallel programming model designed for high scalability and resiliency
- Pioneered by Google. Also designed Percolator – *incrementally processing updates*
 - *Implementation in C++*
- Popularised by the open-source Hadoop project
 - Used at Yahoo!, Facebook, Amazon, ...

Used / has been used for

- At Google
 - Index construction for Google Search, Article clustering for Google News, Statistical machine translation
- At Yahoo!
 - “Web map” powering Yahoo! Search, Spam detection for Yahoo! Mail
- At Facebook
 - Ad optimisation, Spam detection
- In research
 - Astronomical image analysis, Bioinformatics, Analysing Wikipedia conflicts, Natural language processing, Particle physics, Ocean climate simulation
- Among others...

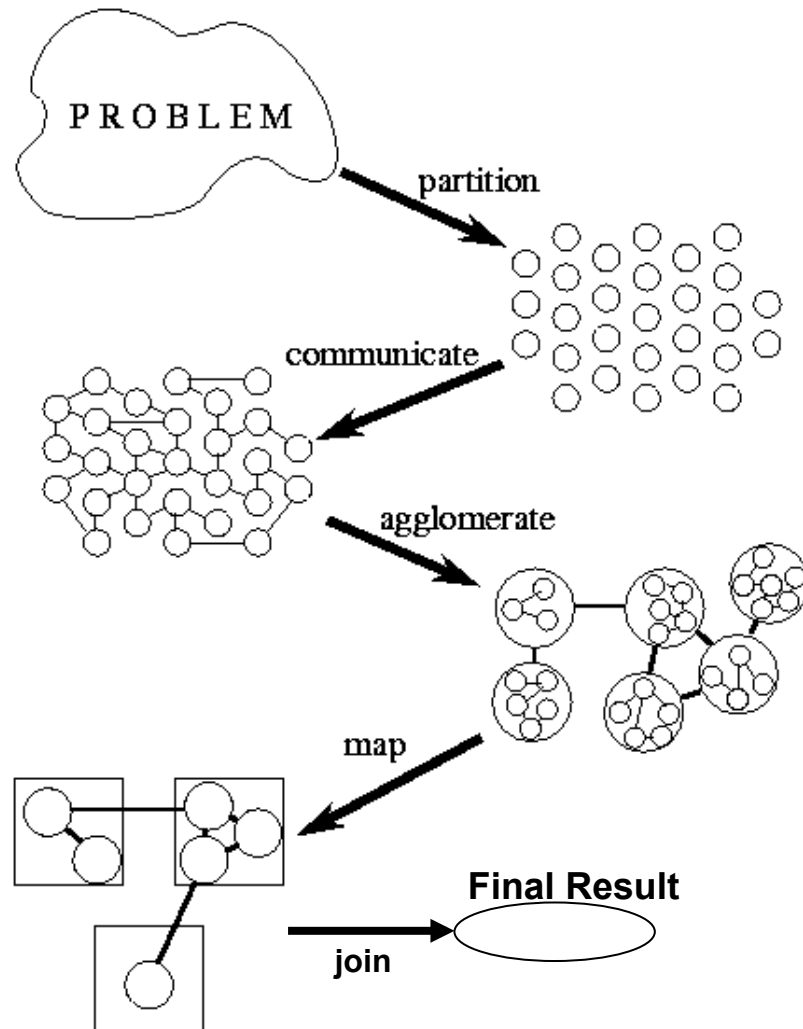
Parallel Programming

- Process of developing programs that express what computations should be executed in parallel.
- Parallel computing
 - use of two or more processors (computers), ***usually within a single system***, working simultaneously to solve a single problem
- Distributed computing
 - any computing that involves ***multiple computers remote from each other*** that each have a role in a computation problem or information processing

Parallel vs. Distributed Computing

<i>Characteristic</i>	<i>Parallel</i>	<i>Distributed</i>
Overall Goal	Speed	Convenience
Interactions	Frequent	Infrequent
Granularity	Fine	Coarse
Reliable	Assumed	Not Assumed

Parallel Programming Methodical Design



● Partitioning

- Computation, data are decomposed into small tasks.

● Communication

- Coordinate task execution

● Agglomeration

- Tasks are combined to larger tasks

● Map

- Assigned to a processor

● Join

Parallel Programming Paradigms

● Data parallelism

- All tasks apply the same set of operations to different data
- Example: for $i \leftarrow 0$ to 99 do
 $a[i] \leftarrow b[i] + c[i]$
 endfor
- Operations may be performed either synchronously or asynchronously
- Fine grained parallelism
- *Grain Size* is the average number of computations performed between communication or synchronization steps

Parallel Programming Paradigms

● Task parallelism

- Independent tasks apply different operations to different data elements

- Example

$a \leftarrow 2$

$b \leftarrow 3$

$m \leftarrow (a + b) / 2$

$s \leftarrow (a^2 + b^2) / 2$

$v \leftarrow s - m^2$

- Concurrent execution of tasks, not statements
- Problem is divided into different tasks
- Tasks are divided between the processors
- Coarse grained parallelism

Traditional HPC systems

- **CPU-intensive computations**

- Relatively small amount of data
- Tightly-coupled applications
- Highly concurrent I/O requirements
- Complex message passing paradigms such as MPI, PVM...
- Developers might need to spend some time designing for failure