

MapReduce

A programming model
for processing large datasets

Part 2

Traditional HPC systems

- **CPU-intensive computations**

- Relatively small amount of data
- Tightly-coupled applications
- Highly concurrent I/O requirements
- Complex message passing paradigms such as MPI, PVM...
- Developers might need to spend some time designing for failure

Challenges

- **Data and storage**
 - Locality, computation close to the data
- **In large-scale systems, nodes fail**
 - MTBF (Mean time between failures) for 1 node = 3 years
 - MTBF for 1000 nodes = 1 day
 - Solution: Built in fault-tolerance
- **Commodity network = low bandwidth**
- **Distributed programming is hard**
 - Solution: simple data-parallel programming model: users structure the application in “map” & “reduce” functions, system distributes data/work and handles faults
 - Not all applications can be parallelised: tightly-coupled computations

What requirements?

- A simple data-parallel programming model, designed for high scalability and resiliency
 - Scalability to large-scale data volumes
 - Automated fault-tolerance at application level rather than relying on high-availability hardware
 - Simplified I/O and tasks monitoring
 - All based on cost-efficient commodity machines (cheap, but unreliable), and commodity network

Core concepts

- Data spread in advance, persistent (in terms of locality), and replicated
- No inter-dependencies / shared nothing architecture
- Applications written in two pieces of code
 - And developers do not have to worry about the underlying issues in networking, jobs interdependencies, scheduling, etc...

The model

- A map function processes a key/value pair to generate a set of intermediate key/value pairs
 - Divides the problem into smaller 'intermediate key/value' pairs
- The reduce function merge all intermediate values associated with the same intermediate key
- Run-time system takes care of:
 - Partitioning the input data across nodes (blocks/chunks typically of 64Mb to 128Mb)
 - Scheduling the data and execution
 - Node failures, replication, re-submissions
 - Coordination among nodes

Map function

- A map function processes a key/value pair to generate a set of intermediate key/value pairs
 - Divides the problem into smaller 'intermediate key/value' pairs
- Map: $(key1, val1) \rightarrow (key2, val2)$
- Ex:
 - $(line-id, text) \rightarrow (word, 1)$
 - $(2, \text{the apple is an apple}) \rightarrow (the, 1), (apple, 1), (is, 1), (an, 1), (apple, 1)$

Reduce function

- The reduce function merge all intermediate values associated with the same intermediate key
- Reduce: $(key2, [val2]) \rightarrow [val3]$
- Ex:
 - $(word, [val_1, val_2, \dots]) \rightarrow (word, \Sigma val_i)$
 - $(the, 1), (apple, 1), (is, 1), (an, 1), (apple, 1) \rightarrow (an, 1), (apple, 2), (is, 1), (the, 1)$

Map/Reduce

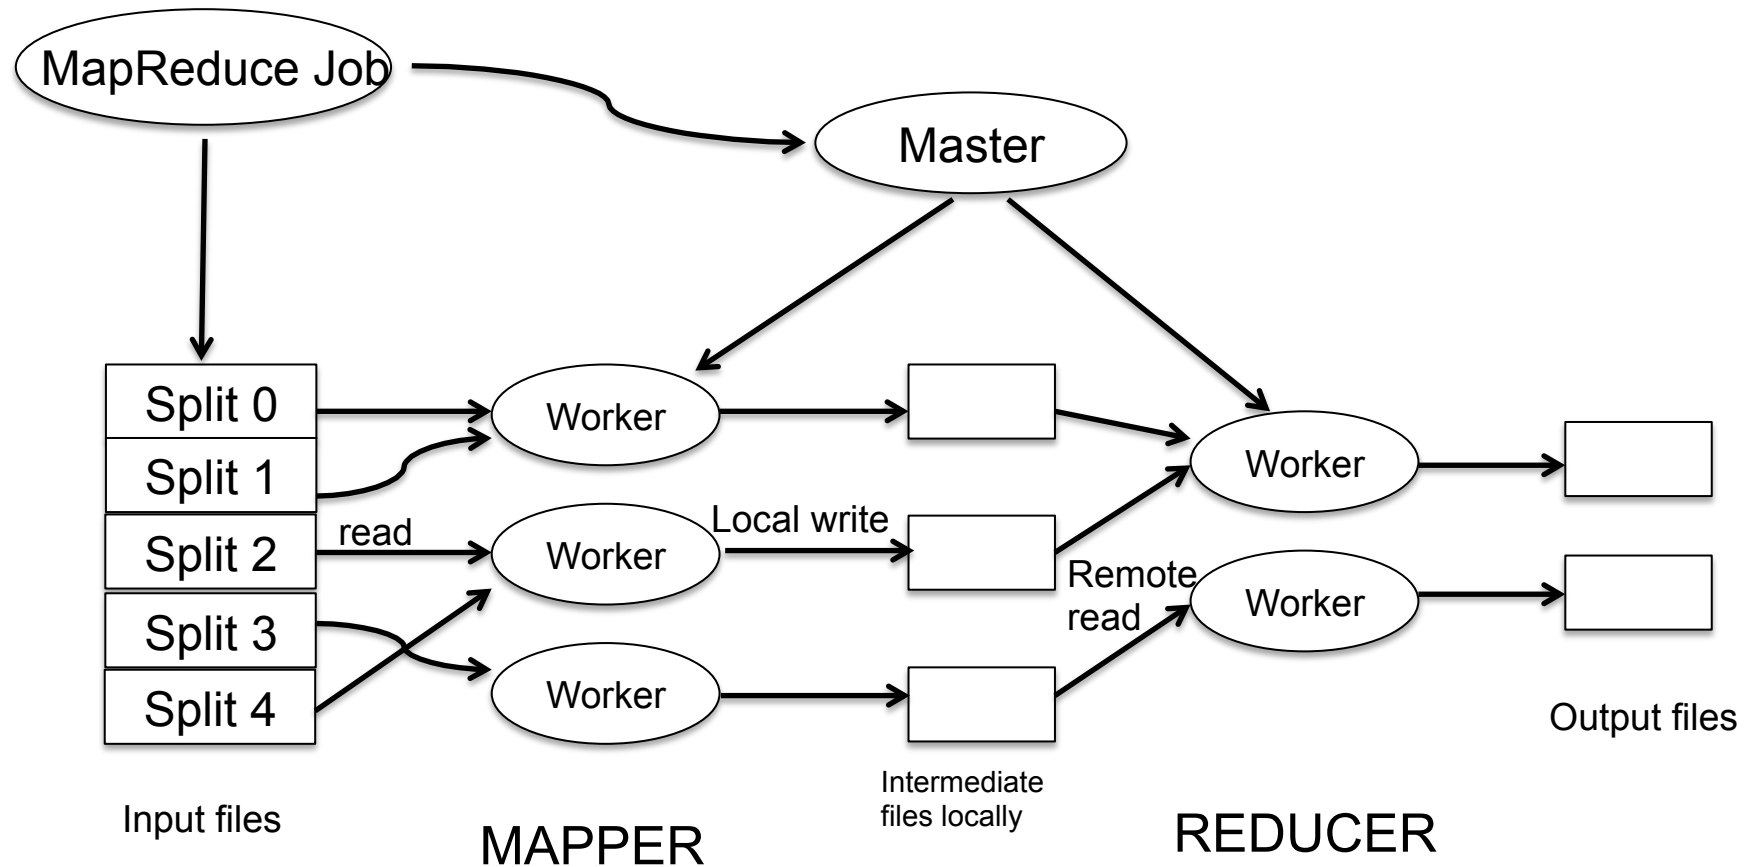
Map

- Extract something of interest from a large number of records

Reduce

- Sort intermediate results
- Aggregate intermediate results
- Generate final output

MapReduce execution



Simple Word Count

```
#key: offset, value: line
```

```
def mapper():
```

```
    for line in open("doc"):
```

```
        for word in line.split():
```

```
            output(word, 1)
```

```
#key: a word, value: iterator over counts
```

```
def reducer():
```

```
    output(key, sum(value))
```